



I tested 4 open-source LLMs against Claude Sonnet 4.6

for enterprise QE automation

On my desktop GPU. Here's what I found.

Swipe →

The experiment



My workstation

- GPU: RTX 5060 Ti — 16 GB VRAM
- CPU: AMD Ryzen 7 9700X (8-core)
- RAM: 64 GB DDR5
- OS: Windows 11 Pro



4 models tested

Devstral Small 2

24B dense · 14 GB

Qwen3.6-35B-A3B

35B MoE (3B active) · 5 GB

Gemma 4 26B

26B MoE (4B active) · 8 GB

Ornith-1.0-9B

9B dense · 5 GB

6-agent pipeline tested end-to-end

Enrichment

Explorer

Builder

Executor

Reviewer

Healer

Scored against Claude Sonnet 4.6 using a custom 9-dimension quality scorecard

The surprise

A 9B model outscored two models 3× its size

Ornith 9B

5 GB VRAM · 10 min



74%

Devstral 24B

14 GB VRAM · 40 min



57%

Gemma 4 26B

8 GB VRAM · 40 min



54%

Claude
90-95%



Why?

Ornith uses self-scaffolding reinforcement learning — it doesn't just generate code, it learns to build an internal framework from your instructions first, then generates code guided by that framework. Bigger models read the same rules but forget them mid-generation.

Same model. Different instructions. Completely different result.



Flat instruction files

5 / 23

steps completed

Fabricated a "Shadow DOM"
explanation and stopped



Skills-based disclosure

24 / 24

steps completed

Full navigation in 10 minutes
Zero hallucinations

HOW you deliver instructions to a local model matters as much as **WHICH**
model you pick.

3 things no benchmark captures

1

Training > Size

Ornith 9B (self-scaffolding RL) scored 74%.
Devstral 24B scored 57%. Gemma 26B scored 54%.
Training methodology beats parameter count
for instruction-following tasks.

2

Instruction architecture is an independent variable

Same model, same scenario.
Flat files: 5/23 steps. Skills-based: 24/24 steps.
No benchmark measures this.
Your framework design is half the battle.

3

Local models hallucinate with confidence

One model fabricated a Shadow DOM explanation.
Another generated locators from a logistics app
that doesn't exist in the target application.
More dangerous than failing silently.

The honest verdict



Not production-ready today at 16 GB VRAM

- Best local Builder score: 74% vs Claude's 90-95%
- Builder accounts for only ~10% of pipeline credits — savings are modest
- 8-9× slower than cloud models
- Hallucination risk requires cloud-based Reviewer as safety net



But the gap is closing — fast

- NVIDIA RTX Spark (Fall 2026): 128 GB unified memory in a laptop
- DGX Spark: 128 GB, runs 70B+ models, ~\$4,000
- Muse Glimmer, Qwen3.6-27B, Laguna XS 2.1 — all shipped since this experiment
- **The question is not IF local models will match cloud. It's WHEN.**

My 9-dimension quality scorecard tells you exactly when that moment arrives.

Run any new model through the same scorecard.
When it scores above 80% with no dimension below 3 — it's ready.



**The full whitepaper,
the framework, and the scorecard
are on my GitHub**

github.com/sreevsr

Srinidhi Sreevatsa

QE Consultant & Architect · 20+ years in Enterprise Quality Engineering

linkedin.com/in/srinidhi-sreevatsa-tnarasipura

#QualityEngineering #LocalLLM #AgenticAI #TestAutomation #Playwright